

CSE 250A. Principles of AI

Probabilistic Reasoning and Decision-Making

Lecture 10 – Learning from complete and incomplete data


Lawrence Saul
Department of Computer Science and Engineering
University of California, San Diego

Fall 2021

Outline

- 1 Review
- 2 Logistic regression
- 3 Learning from incomplete data
- 4 Auxiliary functions

Linear regression



Suppose $Y \in \mathbb{R}$ is a real-valued random variable.
 Then we can use a Gaussian conditional distribution:

$$P(y|\vec{x}) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ -\frac{(y - \vec{w} \cdot \vec{x})^2}{2\sigma^2} \right\}$$

How to learn the weight vector \vec{w} from complete (IID) data?
 This is the problem of **linear regression**.

ML estimation for linear regression

- Log-conditional likelihood

$$\mathcal{L}(\vec{w}) = -\frac{1}{2} \sum_{t=1}^T \left[\log(2\pi\sigma^2) + \frac{(y_t - \vec{w} \cdot \vec{x}_t)^2}{\sigma^2} \right]$$

- Least-squares solution

$$\left. \begin{aligned} \mathbf{A} &= \sum_t \vec{x}_t \vec{x}_t^\top \\ \vec{b} &= \sum_t y_t \vec{x}_t \end{aligned} \right\} \quad \vec{b} = \mathbf{A} \vec{w} \quad \Rightarrow \quad \boxed{\vec{w}_{\text{ML}} = \mathbf{A}^{-1} \vec{b}}$$

- Failure modes

Ill-conditioned problems arise when the inputs $\{\vec{x}_1, \vec{x}_2, \dots, \vec{x}_T\}$ lie in (or very nearly in) a proper subspace of \mathbb{R}^d .

A detour on numerical optimization

How to maximize a multivariable function $f(\vec{\theta})$ over $\vec{\theta} \in \mathbb{R}^d$?

① Analytically — when it is possible

Compute the gradient and solve for where it vanishes:

$$\nabla f = \left(\frac{\partial f}{\partial \theta_1}, \frac{\partial f}{\partial \theta_2}, \dots, \frac{\partial f}{\partial \theta_d} \right) = (0, 0, \dots, 0)$$

② Numerically — via hillclimbing

Perform an iterative local search for a local or global maximum of $f(\vec{\theta})$.

Gradient ascent

- **Iterative update**

$$\vec{\theta} \leftarrow \vec{\theta} + \eta \left(\frac{\partial f}{\partial \vec{\theta}} \right)$$

- **Pros**

Applies to any once-differentiable function.

Converges asymptotically for sufficiently small η .

- **Cons**

Sometimes tricky in practice to tune the learning rate.

No guarantee of monotonic convergence.

No guarantee of global optimality.

Newton's method

- **Iterative update**

$$\vec{\theta} \leftarrow \vec{\theta} - \mathbf{H}^{-1} \left(\frac{\partial f}{\partial \vec{\theta}} \right)$$

- **Pros**

Applies to any twice-differentiable function.

Converges rapidly (when it converges).

Avoids the difficulty of tuning of a learning rate.

- **Cons**

Expensive to compute Hessian matrix $O(d^2)$.

Expensive to solve linear system $O(d^2)$.


Unpredictable and/or unstable when initial estimate is poor.

No guarantee of global optimality.

Outline

- 1 Review
- 2 **Logistic regression**
- 3 Learning from incomplete data
- 4 Auxiliary functions

Logistic regression



Suppose $Y \in \{0, 1\}$ is a binary random variable.
Then we can use a **sigmoid conditional distribution**:

$$P(Y=1|\vec{x}) = \sigma(\vec{w} \cdot \vec{x}) = \frac{1}{1 + e^{-\vec{w} \cdot \vec{x}}}$$

How to learn the parameter $\vec{w} \in \mathbb{R}^d$ from complete data?
This is the problem of **logistic regression**.

Preliminaries

- **Properties of sigmoid function**

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

bounded between 0 and 1

$$\sigma(-z) = 1 - \sigma(z)$$

reflection symmetry

$$\frac{d}{dz}\sigma(z) = \sigma(z)\sigma(-z)$$

derivative

- **IID data**

As usual, we assume a complete data set of IID examples $\{(\vec{x}_t, y_t)\}_{t=1}^T$ where $\vec{x}_t \in \mathbb{R}^d$ and $y_t \in \{0, 1\}$.

Log-conditional likelihood

$$\mathcal{L}(\vec{w}) = \log P(y_1, y_2, \dots, y_T | \vec{x}_1, \vec{x}_2, \dots, \vec{x}_T)$$

$$= \log \prod_{t=1}^T P(y_t | \vec{x}_t)$$

data is IID

$$= \sum_{t=1}^T \log P(y_t | \vec{x}_t)$$

$\log ab = \log a + \log b$

$$= \sum_{t=1}^T \log \left[\sigma(\vec{w} \cdot \vec{x}_t)^{y_t} (1 - \sigma(\vec{w} \cdot \vec{x}_t))^{1-y_t} \right]$$

$y_t \in \{0, 1\}$

$$= \sum_{t=1}^T [y_t \log \sigma(\vec{w} \cdot \vec{x}_t) + (1 - y_t) \log(1 - \sigma(\vec{w} \cdot \vec{x}_t))]$$

$\log a^b = b \log a$

$$= \sum_{t=1}^T [y_t \log \sigma(\vec{w} \cdot \vec{x}_t) + (1 - y_t) \log \sigma(-\vec{w} \cdot \vec{x}_t)]$$

$\sigma(-z) = 1 - \sigma(z)$

Computing the partial derivatives

$$\begin{aligned}
 \mathcal{L}(\vec{w}) &= \sum_t \left[y_t \log \sigma(\vec{w} \cdot \vec{x}_t) + (1-y_t) \log \sigma(-\vec{w} \cdot \vec{x}_t) \right] \\
 \frac{\partial \mathcal{L}}{\partial w_\alpha} &= \sum_t \left[y_t \frac{1}{\sigma(\vec{w} \cdot \vec{x}_t)} \sigma(\vec{w} \cdot \vec{x}_t) \sigma(-\vec{w} \cdot \vec{x}_t) x_{\alpha t} \right. \\
 &\quad \left. + (1-y_t) \frac{1}{\sigma(-\vec{w} \cdot \vec{x}_t)} \sigma(-\vec{w} \cdot \vec{x}_t) \sigma(\vec{w} \cdot \vec{x}_t) (-x_{\alpha t}) \right] \\
 &= \sum_t x_{\alpha t} \left[y_t \sigma(-\vec{w} \cdot \vec{x}_t) - (1-y_t) \sigma(\vec{w} \cdot \vec{x}_t) \right] \\
 &= \sum_t x_{\alpha t} \left[y_t (1 - \sigma(\vec{w} \cdot \vec{x}_t)) - (1-y_t) \sigma(\vec{w} \cdot \vec{x}_t) \right] \\
 &= \sum_t x_{\alpha t} \left[y_t - \sigma(\vec{w} \cdot \vec{x}_t) \right]
 \end{aligned}$$

Interpreting the partial derivatives

- Partial derivative

$$\frac{\partial \mathcal{L}}{\partial w_{\alpha}} = \sum_{t=1}^T x_{\alpha t} \underbrace{\left[\textcolor{blue}{y_t} - \textcolor{red}{\sigma(\vec{w} \cdot \vec{x}_t)} \right]}_{\text{error signal}}$$

- Error signals

For each example (\vec{x}_t, y_t) , the signal compares what the model should predict versus what it does:

$$\begin{array}{rcccl} \text{error signal} & = & \text{target label} & - & \text{model prediction} \\ & & \textcolor{blue}{y_t} & & \textcolor{red}{\sigma(\vec{w} \cdot \vec{x}_t)} \\ & & P(Y=1|Y=y_t) & & P(Y=1|X=\vec{x}_t) \end{array}$$

Maximizing the log-conditional likelihood

- Where does the gradient vanish?

$$\frac{\partial \mathcal{L}}{\partial w_{\alpha}} = 0 \quad \implies \quad \sum_{t=1}^T x_{\alpha t} y_t = \sum_{t=1}^T x_{\alpha t} \sigma(\vec{w} \cdot \vec{x}_t)$$

- The good news:

We have d equations, one for each $\alpha \in \{1, 2, \dots, d\}$.

And we have d unknowns $\vec{w} = (w_1, w_2, \dots, w_d)$.

- The bad news:

These equations are nonlinear — because $\sigma(z)$ is nonlinear.

There is no way to solve them in closed form.

Maximizing the log-conditional likelihood

*If we can't do it analytically, then we must do it numerically.
What are the simplest hillclimbing methods for this problem?*

1 Gradient ascent

$$\vec{w} \leftarrow \vec{w} + \eta \left(\frac{\partial \mathcal{L}}{\partial \vec{w}} \right)$$

2 Newton's method

$$\vec{w} \leftarrow \vec{w} - \mathbf{H}^{-1} \left(\frac{\partial \mathcal{L}}{\partial \vec{w}} \right)$$

Partial derivatives

- First partial derivatives

$$\frac{\partial \mathcal{L}}{\partial w_{\alpha}} = \sum_{t=1}^T x_{\alpha t} \left[y_t - \sigma(\vec{w} \cdot \vec{x}_t) \right]$$

- Second partial derivatives

$$\begin{aligned} \frac{\partial^2 \mathcal{L}}{\partial w_{\alpha} \partial w_{\beta}} &= \frac{\partial}{\partial w_{\beta}} \left(\sum_{t=1}^T x_{\alpha t} \left[y_t - \sigma(\vec{w} \cdot \vec{x}_t) \right] \right) \\ &= - \sum_{t=1}^T x_{\alpha t} \left[\sigma(\vec{w} \cdot \vec{x}_t) \sigma(-\vec{w} \cdot \vec{x}_t) x_{\beta t} \right] \\ &= - \sum_{t=1}^T \sigma(\vec{w} \cdot \vec{x}_t) \sigma(-\vec{w} \cdot \vec{x}_t) x_{\alpha t} x_{\beta t} \end{aligned}$$

Gradient and Hessian in matrix-vector notation

- **Gradient**

$$\nabla \mathcal{L} = \frac{\partial \mathcal{L}}{\partial \vec{w}} = \sum_{t=1}^T \left[y_t - \sigma(\vec{w} \cdot \vec{x}_t) \right] \vec{x}_t$$


- **Hessian**

$$\mathbf{H} = \frac{\partial^2 \mathcal{L}}{\partial \vec{w} \partial \vec{w}^\top} = - \sum_{t=1}^T \sigma(\vec{w} \cdot \vec{x}_t) \sigma(-\vec{w} \cdot \vec{x}_t) \vec{x}_t \vec{x}_t^\top$$


You will need to compute these for HW 5. Questions?

Logistic regression in HW 5

How well can this model distinguish images of **3s** versus **5s**?



$y = 0$



$y = 1$


Recommendations:

- Set $\eta = \frac{0.2}{T}$ for gradient ascent.
- Initialize $\vec{w} = (0, 0, \dots, 0)$ for Newton's method.

Global optimality

- **Theorem**

*The log-conditional likelihood $\mathcal{L}(\vec{w})$ for logistic regression is a **concave** function of \vec{w} .*



- **Corollary**

$\mathcal{L}(\vec{w})$ has no spurious local maxima.

You should all converge to the same solution for HW 5!

- **Proof sketch**

The Hessian is negative semidefinite: $\vec{v}^\top \mathbf{H} \vec{v} \leq 0$ for all $\vec{v} \in \mathbb{R}^d$.
This is a sufficient condition for concavity.

Outline

- 1 Review
- 2 Logistic regression
- 3 **Learning from incomplete data**
- 4 Auxiliary functions


Learning from incomplete data with tabular CPTs

ASSUMPTIONS

- 1 The DAG is fixed (and known) over a finite set of discrete random variables $\{X_1, X_2, \dots, X_n\}$.
- 2 CPTs enumerate $P(X_i = x | \text{pa}(X_i) = \pi)$ as lookup tables; each must be estimated for all values of x and π .
- 3 The data is IID, but only consists of T **partially** complete instantiations of the nodes in the BN.

Toy example

- Fixed DAG over binary random variables



$$X_1 \in \{0, 1\}$$

$$X_2 \in \{0, 1\}$$

$$X_3 \in \{0, 1\}$$


$$X_4 \in \{0, 1\}$$

- Incomplete data set

example	X_1	X_2	X_3	X_4
1	1	?	0	1
2	0	1	?	0
3	?	?	?	1
:	:	:	:	:
T	?	1	1	0

How to choose the CPTs so that the BN maximizes the probability of this data set?

A more interesting example ...



How to build a movie recommendation system?

- Collect a data set of movie ratings:


$$\begin{bmatrix} + & - & + & - & ? & ? & + \\ - & ? & ? & + & + & ? & ? \\ + & + & + & + & + & + & + \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \\ - & - & - & - & - & ? & - \\ ? & ? & + & ? & ? & ? & - \end{bmatrix}$$

+	liked
-	disliked
?	not seen

(user-item matrix)

- Build a model of user profiles and fill in the missing ratings.
But what model to build? (HW 8)

Naive Bayes model with incomplete data



- Movie recommender system

$Z \in \{1, 2, \dots, k\}$ type of movie-goer
 $R_i \in \{0, 1\}$ rating for i^{th} movie

- Incomplete data set

student	Z	R_1	R_2	R_3	R_4	\dots
1	?	0	1	1	?	\dots
2	?	1	?	0	1	\dots
3	?	0	0	?	1	\dots
\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots
T	?	?	1	0	?	\dots

Note that the variable Z is **never observed**.


Learning from incomplete data

• Notation

H_t = set of hidden (latent) variables for t^{th} example

V_t = set of visible (observed) variables for t^{th} example

• Illustration



example	X_1	X_2	X_3	X_4
1	1	?	0	1
2	0	1	?	0
3	?	?	?	1
:	:	:	:	:

$$H_1 = \{X_2\}$$

$$H_2 = \{X_3\}$$

$$H_3 = \{X_1, X_2, X_3\}$$

$$V_1 = \{X_1, X_3, X_4\}$$

$$V_2 = \{X_1, X_2, X_4\}$$

$$V_3 = \{X_4\}$$

Computing the log-likelihood with **incomplete** data

$$\mathcal{L} = \log P(\mathbf{data})$$

$$= \log \prod_{t=1}^T P(V_t = v_t)$$

data is IID

$$= \sum_{t=1}^T \log P(V_t = v_t)$$

$\log ab = \log a + \log b$

$$= \sum_{t=1}^T \log \sum_h P(H_t = h, V_t = v_t)$$

marginalization

$$= \sum_{t=1}^T \log \sum_h P(X_1 = x_1, X_2 = x_2, \dots, X_n = x_n) \Big|_{\{H_t = h, V_t = v_t\}}$$

joint

$$= \sum_{t=1}^T \log \sum_h \prod_{i=1}^n P(X_i = x_i | \text{pa}_i = \pi_i) \Big|_{\{H_t = h, V_t = v_t\}}$$

product rule

Complete versus incomplete data

- **Complete data**

$$\mathcal{L} = \sum_{i, \pi, x} \text{count}(X_i = x, \text{pa}_i = \pi) \log P(X_i = x | \text{pa}_i = \pi)$$

The CPTs at different nodes are decoupled!

We can compute ML estimates in closed form.

- **Incomplete data**

$$\mathcal{L} = \sum_{t=1}^T \log \sum_{\textcolor{red}{h}} \prod_{i=1}^n P(X_i = x_i | \text{pa}_i = \pi_i) \Big|_{\{H_t = h, V_t = v_t\}}$$

The CPTs are potentially all coupled.

How to proceed?

Outline

- 1 Review
- 2 Logistic regression
- 3 Learning from incomplete data
- 4 **Auxiliary functions**

How to maximize $f(\vec{\theta})$?

1 Gradient ascent

$$\vec{\theta} \leftarrow \vec{\theta} + \eta \left(\frac{\partial f}{\partial \vec{\theta}} \right)$$

- ✗ Tedious to tune η ?
- ✗ Not monotonically convergent.

2 Newton's method

$$\vec{\theta} \leftarrow \vec{\theta} - \mathbf{H}^{-1} \left(\frac{\partial f}{\partial \vec{\theta}} \right)$$

- ✗ Expensive for large problems.
- ✗ Fast but unstable.

3 Auxiliary function

$$\vec{\theta}_{\text{new}} = \underset{\vec{\theta}}{\operatorname{argmax}} Q(\vec{\theta}, \vec{\theta}_{\text{old}})$$

- ✓ No learning rate.
- ✓ Monotonically convergent.

Auxiliary functions

- Definition

A function $Q(\vec{\theta}', \vec{\theta})$ is called an *auxiliary function* for the *objective function* $f(\vec{\theta})$ if it satisfies two properties:

(i) $Q(\vec{\theta}, \vec{\theta}) = f(\vec{\theta})$ for all $\vec{\theta}$

equality

(ii) $Q(\vec{\theta}', \vec{\theta}) \leq f(\vec{\theta}')$ for all $\vec{\theta}, \vec{\theta}'$

lower bound

- Theorem

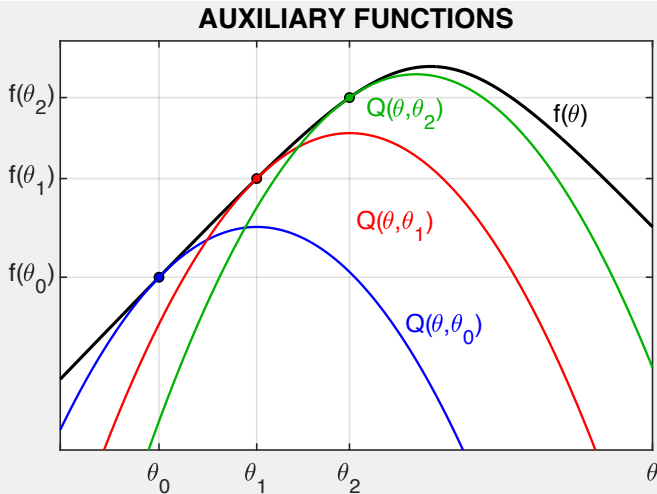
Let $Q(\vec{\theta}', \vec{\theta})$ be an auxiliary function for the objective function $f(\vec{\theta})$. Then the update rule

$$\vec{\theta}_{\text{new}} = \operatorname{argmax}_{\vec{\theta}} Q(\vec{\theta}, \vec{\theta}_{\text{old}})$$



converges monotonically to a stationary point with

$$f(\vec{\theta}_{\text{new}}) \geq f(\vec{\theta}_{\text{old}}).$$

Visualization



Versus Newton's method



The quadratic approximation in Newton's method is **not** guaranteed to provide a lower bound on the objective function.

Proof of monotonic convergence

• Proof

Consider the update rule $\vec{\theta}_{\text{new}} = \operatorname{argmax}_{\vec{\theta}} Q(\vec{\theta}, \vec{\theta}_{\text{old}})$.
 Then we have

$$\begin{aligned}
 f(\vec{\theta}_{\text{new}}) &\geq Q(\vec{\theta}_{\text{new}}, \vec{\theta}_{\text{old}}) && \boxed{\text{property (ii)}} \\
 &\geq Q(\vec{\theta}_{\text{old}}, \vec{\theta}_{\text{old}}) && \boxed{\text{argmax update}} \\
 &= f(\vec{\theta}_{\text{old}}) && \boxed{\text{property (i)}}
 \end{aligned}$$

Iterating this process, we have:

$$f(\vec{\theta}_0) \leq f(\vec{\theta}_1) \leq f(\vec{\theta}_2) \leq \dots \leq f(\vec{\theta}_n).$$

Next lecture: ML estimation for incomplete data!